Rohan Varshney

rohanvar@gmail.com | rohanvarshney.github.io | linkedin.com/in/rohanvarshney | github.com/rohanvarshney | US Citizen

Experience

Lyft *New York City, NY*

Senior Software Engineer (Features & Metrics) *December 2025 – Present*

- Exercised technical ownership of architectural strategy to scale Lyft's ML Feature Store for European market expansion.
- Modeled metadata-driven access control framework to enforce GDPR compliance and PII protection at individual feature level, successfully unblocking mission-critical fraud detection services for launch.
- Reduced AWS DynamoDB provisioning costs by $1M+/year by deprecating redundant GSIs across feature store backing tables.

Senior Software Engineer (Streaming Compute) ***Blog*** | ***Video*** *May 2023 – December 2025*

- Implemented OpenSearch vector search support for embeddings, unlocking LLM & RAG use cases for ML Platform.
- Led $450K+/yr of Infra cost-saving optimizations via Astronomer migration, Redis/EKS tuning, and backing store unification.
- Evaluated managed solution for Feature Service, authoring 200+ pages of documents detailing current system capabilities, use cases, customers, hypothetical migration plans, & evaluation criteria. Led meetings with Director, Staffs, & TPMs.
- Spearheaded a cross-functional Data Contract initiative for Tier 1/2 feature snapshot tables, implementing automated SLA monitoring and PagerDuty alerting to mitigate data quality risks and fiscal repercussions of stale data.
- Deployed new Offline Data Python SDK that standardized feature discovery & selection and dataset generation workflows, saving ML Engineers across the company ∼5 days of development time per model.
- Awarded panding patent **(US20260220581)** for **"Driver Referral Letters"**; conceived and engineered the initial MVP during company hackathon, was awarded overall runner-up and subsequently greenlit for full-scale production product launch.

Software Engineer II (Data Persistence) *August 2022 – May 2023*

- Automated T0 data consolidation workflows, saving hundreds of LUS engineering hours across 20+ table onboardings.
- Cut team SLA breaches, alerts, and failed redrives by 75% via systematic automation and stability improvements.
- Orchestrated work of 6 Data Platform contractors to migrate and deprecate the Hive query engine for Lyft's largest DAGs & operators, personally unblocking complex use cases regarding Spark syntax and rollouts.

Software Engineer I (Data Fabric & Data IDL) *July 2021 – August 2022*

- Pioneered new on-the-fly (un)marshal library functionality to drive the migration from cJSON to Google Protobuf for all 450+ microservices and 2400+ analytics events, saving ∼15% CPU fleet-wide.
- Optimized analytics event proxy sidecars by leveraging a native schema-retrieval API, reducing EC2 provisioning costs and memory footprint by up to ∼80% for all event-emitting services at Lyft.

Microsoft *Redmond, WA*

Software Engineering Intern *May 2020 – August 2020*

- Scaled internal Microsoft Teams chatbot via 8 new deployment features and automated Azure Storage FAQs, boosting agent utilization and cutting on-call TTI/TTR by 4-6 hours per request.

Lyft *San Francisco, CA*

Software Engineering Intern *January 2020 – May 2020*

- Engineered heatmaps using GeoPandas, WKT geometries, and Geohash weights to visualize aggregate regional quality metrics.
- Built 10+ Mode dashboards enabling validation of metric heuristic iterations and OpenStreetMap (OSM) map changes.

Amazon.com *Seattle, WA*

Software Development Engineering Intern *May 2019 – August 2019*

- Developed autocomplete & bookmark migration for AWS AEA (enterprise access) Android app, improved CX for 5,500+ users.
- Automated manual dependencies and dynamic service retrieval to reduce microservice onboarding from 1 month to sub-hour.

Manhattan Associates *Atlanta, GA*

Software Engineering R&D Co-Op *January 2019 – May 2019*

Raytheon *Indianapolis, IN*

Software Engineering Intern *May 2018 – August 2018*

Education

Georgia Institute of Technology *Atlanta, GA*

B.S. in Computer Science (Modeling / Simulation & Intelligence Concentrations)

Skills

Languages: Python, Golang, Java, Kotlin, SQL (Hive, Spark, Trino), TypeScript, Bash, C#, C

Technologies: Jupyter/Databricks Notebooks, Amazon Web Services (EKS, S3, Kinesis, Kafka, DynamoDB, Redis, OpenSearch), Docker, Apache Flink & Beam, NumPy & Pandas, Postman, Angular, Node.js, Cursor, GenAI/LLM (RAG)