

ROHAN VARSHNEY

rohanvar@gmail.com | rohanvarshney.github.io | linkedin.com/in/rohanvarshney | github.com/rohanvarshney | US Citizen

EDUCATION

Georgia Institute of Technology

Atlanta, GA

B.S. in Computer Science (Modeling / Simulation & Intelligence Concentrations)

EXPERIENCE

Lyft

New York City, NY

Senior Software Engineer (Features & Metrics)

December 2025 – Present

- Exercised technical ownership of strategy to scale Lyft's ML Feature Store for European market expansion.
- Engineered metadata-driven access control framework to enforce GDPR compliance and PII protection at individual feature level, successfully unblocking mission-critical fraud detection services.

Senior Software Engineer (Streaming Compute) [Blog](#) | [Video](#) ↗

May 2023 – December 2025

- Engineered OpenSearch ANN/KNN vector search support for Embeddings, unlocking LLM & RAG use cases for ML Platform.
- Led \$450K+/yr of cost-saving optimizations via orchestration tool migration, Redis/EKS tuning, and backing store unification.
- Evaluated managed solution for Feature Service, authoring 200+ pages of documents detailing current system capabilities, use cases, customers, hypothetical migration plans, & evaluation criteria. Led meetings with Director, Staffs, & TPMs.
- Spearheaded a cross-functional Data Contract initiative for T1/T2 feature snapshot tables, implementing automated SLA monitoring and PagerDuty alerting to mitigate quality risks.
- Architected and deployed new Offline Data Python SDK that standardized feature selection and dataset generation workflows, saving ML Engineers ~5 days of development time per model.

Software Engineer II (Data Persistence)

August 2022 – May 2023

- Automated T0 data consolidation workflows, saving hundreds of engineering hours across 20+ table onboardings.
- Cut SLA breaches, alerts, and failed redrives by 75% via systematic automation and stability improvements.
- Orchestrated the work of 6 Data Platform contractors to migrate and deprecate the Hive query engine for Lyft's largest DAGs, personally unblocking complex use cases regarding Spark syntax and rollouts.

Software Engineer I (Data Fabric & Data IDL)

July 2021 – August 2022

- Pioneered new on-the-fly (un)marshal library functionality to drive the migration from cJSON to Google Protobuf for 450+ services and 2400+ analytics events, saving ~15% CPU fleet-wide.
- Optimized analytics event proxy sidecars by leveraging a native schema-retrieval API, reducing EC2 provisioning costs and memory footprint by up to ~80% for all event-emitting services.

Microsoft

Redmond, WA

Software Engineering Intern

May 2020 – August 2020

- Scaled internal Microsoft Teams chatbot via 8 new deployment features and automated Azure Storage FAQs, boosting utilization and cutting on-call TTI/TTR by 4-6 hours per request.

Lyft

San Francisco, CA

Software Engineering Intern

January 2020 – May 2020

- Engineered heatmaps using GeoPandas, WKT geometries, and Geohash weights to visualize aggregate regional quality metrics.
- Developed 10+ Mode dashboards enabling validation of metric heuristic iterations and OpenStreetMap (OSM) map changes.

Amazon.com

Seattle, WA

Software Development Engineering Intern

May 2019 – August 2019

- Developed autocomplete and bookmark migration for AWS AEA Android app, improving CX for 5,500+ users accessing corporate mobile services.
- Automated manual dependencies and dynamic service retrieval via AWS (Lambda, DynamoDB, API Gateway, SNS, S3), reducing onboarding from 1 month to sub-hour.

Manhattan Associates

Atlanta, GA

Software Engineering R&D Co-Op

January 2019 – May 2019

Raytheon

Indianapolis, IN

Software Engineering Intern

May 2018 – August 2018

SKILLS

Languages: Python, Golang, Java, Kotlin, SQL (Hive, Spark, Trino), TypeScript, Bash, C#, C

Technologies: Jupyter/Databricks Notebooks, Amazon Web Services (EKS, S3, Kinesis, Kafka, DynamoDB, Redis, OpenSearch), Docker, Apache Flink & Beam, NumPy & Pandas, Postman, Angular, Node.js, Cursor, GenAI/LLM (RAG)